

MarchMadness

Sam Porter

5/11/2020

```
# Introduction
## This project will be looking at College Basketball data from the 2014-2019 seasons.
## The data was collected from the following location: https://www.kaggle.com/andrewsundberg/college-ba
## The goal of this project will be to predict if a team will make it to the March Madness tournament,
## In order to determine if a user will make it to the March Madness Tournament, we will need to perform
## We will predict based off a binary variable (1 = you made the March Madness Tournament and 0 = you d
## Given this is a binary variable, we will clean up the data and perform a logistic regression, decisi
## To assess the model accuracy, we will review a ROC curve and confusion matrix to determine the overa

### Install Packages
install.packages('dplyr', repos = "http://cran.us.r-project.org")

## package 'dplyr' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\portesa\AppData\Local\Temp\RtmpAx3mk0\downloaded_packages

install.packages('caret', repos = "http://cran.us.r-project.org")

## package 'caret' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\portesa\AppData\Local\Temp\RtmpAx3mk0\downloaded_packages

install.packages('purrr', repos = "http://cran.us.r-project.org")

## package 'purrr' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\portesa\AppData\Local\Temp\RtmpAx3mk0\downloaded_packages

install.packages('tidyr', repos = "http://cran.us.r-project.org")

## package 'tidyr' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\portesa\AppData\Local\Temp\RtmpAx3mk0\downloaded_packages
```

```
install.packages('ggplot2', repos = "http://cran.us.r-project.org")
```

```
## package 'ggplot2' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\portesa\AppData\Local\Temp\RtmpAx3mk0\downloaded_packages
```

```
install.packages('InformationValue', repos = "http://cran.us.r-project.org")
```

```
## package 'InformationValue' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\portesa\AppData\Local\Temp\RtmpAx3mk0\downloaded_packages
```

```
### Load necessary packages
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(purrr)
```

```
##
## Attaching package: 'purrr'

## The following object is masked from 'package:caret':
##
##   lift
```

```
library(tidyr)
library(ggplot2)
library(InformationValue)
```

```
##
## Attaching package: 'InformationValue'
```

```
## The following objects are masked from 'package:caret':
##
##     confusionMatrix, precision, sensitivity, specificity
```

```
library(rpart)
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
##
##     margin
```

```
## The following object is masked from 'package:dplyr':
##
##     combine
```

```
# Data Import
### Import 5 years worth of datasets
data1 <- read.csv("C:\\Users\\portesa\\Desktop\\Dataset\\cbb14.csv")
data1 <- subset(data1, select=-c(REC))
data2 <- read.csv("C:\\Users\\portesa\\Desktop\\Dataset\\cbb15.csv")
data2 <- subset(data2, select=-c(POSTSEASON))
data3 <- read.csv("C:\\Users\\portesa\\Desktop\\Dataset\\cbb16.csv")
data3 <- subset(data3, select=-c(POSTSEASON))
data4 <- read.csv("C:\\Users\\portesa\\Desktop\\Dataset\\cbb17.csv")
data4 <- subset(data4, select=-c(POSTSEASON))
data5 <- read.csv("C:\\Users\\portesa\\Desktop\\Dataset\\cbb18.csv")
data5 <- subset(data5, select=-c(POSTSEASON))
test_data <- read.csv("C:\\Users\\portesa\\Desktop\\Dataset\\cbb19.csv")
test_data <- subset(test_data, select=-c(POSTSEASON))
```

```
### Combine Data from 2014-2018 to serve as our training set
data <- union(data1,data2)
data <- union(data,data3)
data <- union(data, data4)
data <- union(data, data5)
```

```
## We have several columns of data. Let me first explain what is being stored in each column.
### TEAM = The Division 1 college basketball school
### CONF = The Athletic Conference in which the school participates in
### G = Number of Games played
### W = Number of games won
### ADJOE = Adjusted Offensive Efficiency (An estimate of the offensive efficiency (points scored per 100 possessions))
### ADJDE = Adjusted Defensive Efficiency (An estimate of the defensive efficiency (points allowed per 100 possessions))
### BARTHAG = Power Rating (Chance of beating an average D1 team)
### EFG_O = Effective Field Goal Percentage Shot
```

```

### EFG_D = Effective Field Goal Percentage Allowed
### TOR = Turnover Percentage Allowed (Turnover Rate)
### TORD = Turnover Percentage Committed (Steal Rate)
### ORB = Offensive Rebound Percentage
### DRB = Defensive Rebound Percentage
### FTR = Free Throw Rate (How often the given team shoots Free Throws)
### FTRD = Free Throw Rate Allowed
### 2P_O = Two-Point Shooting Percentage
### 2P_D = Two-Point Shooting Percentage Allowed
### 3P_O = Three-Point Shooting Percentage
### 3P_D = Three-Point Shooting Percentage Allowed
### ADJ_T = Adjusted Tempo (An estimate of the temp (possessions per 40 minutes) a team would have against the bubble)
### WAB = Wins above Bubble (The bubble refers to the cut off between making the NCC March Madness Tournament and not making it)
### POSTSEASON = Round where the given team was eliminated or where their season ended
### Seed = Seed in the NCAA March Madness Tournament

```

```
head(data)
```

```

##          TEAM CONF  G  W ADJOE ADJDE BARTHAG EFG_O EFG_D  TOR TORD  ORB  DRB
## 1 Louisville Amer 37 31 118.8  87.6  0.9710  53.5  43.9 15.3 25.0 37.1 32.7
## 2  Arizona P12 38 33 116.2  87.4  0.9636  51.7  42.3 15.7 19.1 36.4 27.3
## 3  Florida SEC 39 36 115.9  88.4  0.9575  52.2  45.4 17.5 21.3 35.3 28.0
## 4  Virginia ACC 37 30 114.6  89.5  0.9449  50.8  44.2 16.5 18.4 33.9 25.8
## 5 Wisconsin B10 38 30 122.7  95.9  0.9441  53.3  47.2 12.7 15.3 28.1 27.4
## 6    Duke ACC 35 26 125.9  98.6  0.9432  53.8  49.3 14.6 18.5 35.2 31.3
##      FTR FTRD X2P_O X2P_D X3P_O X3P_D ADJ_T  WAB SEED
## 1 41.2 38.4  52.7  44.3  36.8  28.6 68.8  5.3   4
## 2 41.0 34.2  50.7  40.2  36.4  32.0 64.3  9.4   1
## 3 42.4 31.2  51.3  43.5  35.9  33.0 63.1 11.7   1
## 4 42.0 32.5  49.0  42.1  36.9  32.3 61.2  8.2   1
## 5 42.7 27.1  51.3  45.9  37.6  34.1 63.9  7.9   2
## 6 38.8 40.8  50.3  50.3  39.5  30.7 66.7  6.5   3

```

```

# Data Exploration/Data Visualization and Data Cleaning
## In this section we will be taking a look at the distributions of the different features.
## This will help give us a better understanding of the data and what type of feature engineering we may need.
## Before moving forward, we should check and address how we will deal with null values, convert any categorical variables to numerical, etc.

## First let's look at a summary of the data and check to see how many NA values are in the dataset by
summary(data)

```

```

##          TEAM          CONF          G          W
## Length:1755      Length:1755      Min.   :15.00      Min.   : 0.00
## Class :character Class :character 1st Qu.:30.00      1st Qu.:11.00
## Mode  :character Mode  :character Median :31.00      Median :16.00
##                                     Mean  :31.45      Mean   :16.23
##                                     3rd Qu.:33.00      3rd Qu.:21.00
##                                     Max.   :40.00      Max.   :38.00
##
##          ADJOE          ADJDE          BARTHAG          EFG_O
## Min.   : 76.7      Min.   : 84.00      Min.   :0.0077      Min.   :39.4
## 1st Qu.: 98.8      1st Qu.: 99.15      1st Qu.:0.2842      1st Qu.:47.8
## Median :103.4      Median :103.80      Median :0.4740      Median :49.8

```

```
## Mean :103.8 Mean :103.79 Mean :0.4938 Mean :49.9
## 3rd Qu.:108.5 3rd Qu.:108.30 3rd Qu.:0.7135 3rd Qu.:51.9
## Max. :129.1 Max. :124.00 Max. :0.9842 Max. :59.8
##
## EFG_D TOR TORD ORB
## Min. :39.60 Min. :11.90 Min. :10.20 Min. :15.00
## 1st Qu.:48.10 1st Qu.:17.20 1st Qu.:17.00 1st Qu.:27.15
## Median :50.10 Median :18.50 Median :18.40 Median :29.90
## Mean :50.09 Mean :18.54 Mean :18.47 Mean :29.86
## 3rd Qu.:52.00 3rd Qu.:19.80 3rd Qu.:19.80 3rd Qu.:32.55
## Max. :59.50 Max. :26.10 Max. :28.00 Max. :42.10
##
## DRB FTR FTRD X2P_0 X2P_D
## Min. :18.40 Min. :21.6 Min. :22.1 Min. :38.30 Min. :37.70
## 1st Qu.:28.00 1st Qu.:32.9 1st Qu.:32.3 1st Qu.:46.60 1st Qu.:46.70
## Median :30.00 Median :36.4 Median :36.5 Median :48.70 Median :49.00
## Mean :30.06 Mean :36.6 Mean :36.9 Mean :48.81 Mean :48.98
## 3rd Qu.:32.00 3rd Qu.:40.2 3rd Qu.:41.0 3rd Qu.:51.00 3rd Qu.:51.30
## Max. :40.40 Max. :58.6 Max. :60.7 Max. :62.60 Max. :59.80
##
## X3P_0 X3P_D ADJ_T WAB
## Min. :25.20 Min. :27.10 Min. :57.20 Min. : -25.200
## 1st Qu.:32.60 1st Qu.:33.10 1st Qu.:65.70 1st Qu.: -12.900
## Median :34.60 Median :34.70 Median :67.90 Median : -8.300
## Mean :34.57 Mean :34.76 Mean :67.91 Mean : -7.768
## 3rd Qu.:36.40 3rd Qu.:36.40 3rd Qu.:70.00 3rd Qu.: -3.050
## Max. :44.10 Max. :43.10 Max. :83.40 Max. : 13.100
##
## SEED
## Min. : 1.000
## 1st Qu.: 5.000
## Median : 9.000
## Mean : 8.794
## 3rd Qu.:13.000
## Max. :16.000
## NA's :1415
```

```
sapply(data, function(x) sum(is.na(x)))
```

```
## TEAM CONF G W ADJOE ADJDE BARTHAG EFG_0 EFG_D TOR
## 0 0 0 0 0 0 0 0 0 0
## TORD ORB DRB FTR FTRD X2P_0 X2P_D X3P_0 X3P_D ADJ_T
## 0 0 0 0 0 0 0 0 0 0
## WAB SEED
## 0 1415
```

```
## View variable types to see which variables need to be converted to categorical
sapply(data,class)
```

```
## TEAM CONF G W ADJOE ADJDE
## "character" "character" "integer" "integer" "numeric" "numeric"
## BARTHAG EFG_0 EFG_D TOR TORD ORB
## "numeric" "numeric" "numeric" "numeric" "numeric" "numeric"
```

```
##           DRB           FTR           FTRD           X2P_O           X2P_D           X3P_O
## "numeric" "numeric" "numeric" "numeric" "numeric" "numeric"
##           X3P_D           ADJ_T           WAB           SEED
## "numeric" "numeric" "numeric" "integer"
```

```
## View how many conferences we have and then convert them to a numeric categorical variable
```

```
data <- data %>% mutate(CONF = toupper(CONF))
conference <- data %>% group_by(CONF) %>% summarize(ct = n())
conference <- conference %>% mutate(CONF_rating = as.numeric(factor(conference$CONF, levels = conference$CONF)))
data <- data %>% inner_join(conference)
```

```
## Joining, by = "CONF"
```

```
## Same as above, but with the test_data
```

```
test_data <- test_data %>% mutate(CONF = toupper(CONF))
conference <- test_data %>% group_by(CONF) %>% summarize(ct = n())
conference <- conference %>% mutate(CONF_rating = as.numeric(factor(conference$CONF, levels = conference$CONF)))
test_data <- test_data %>% inner_join(conference)
```

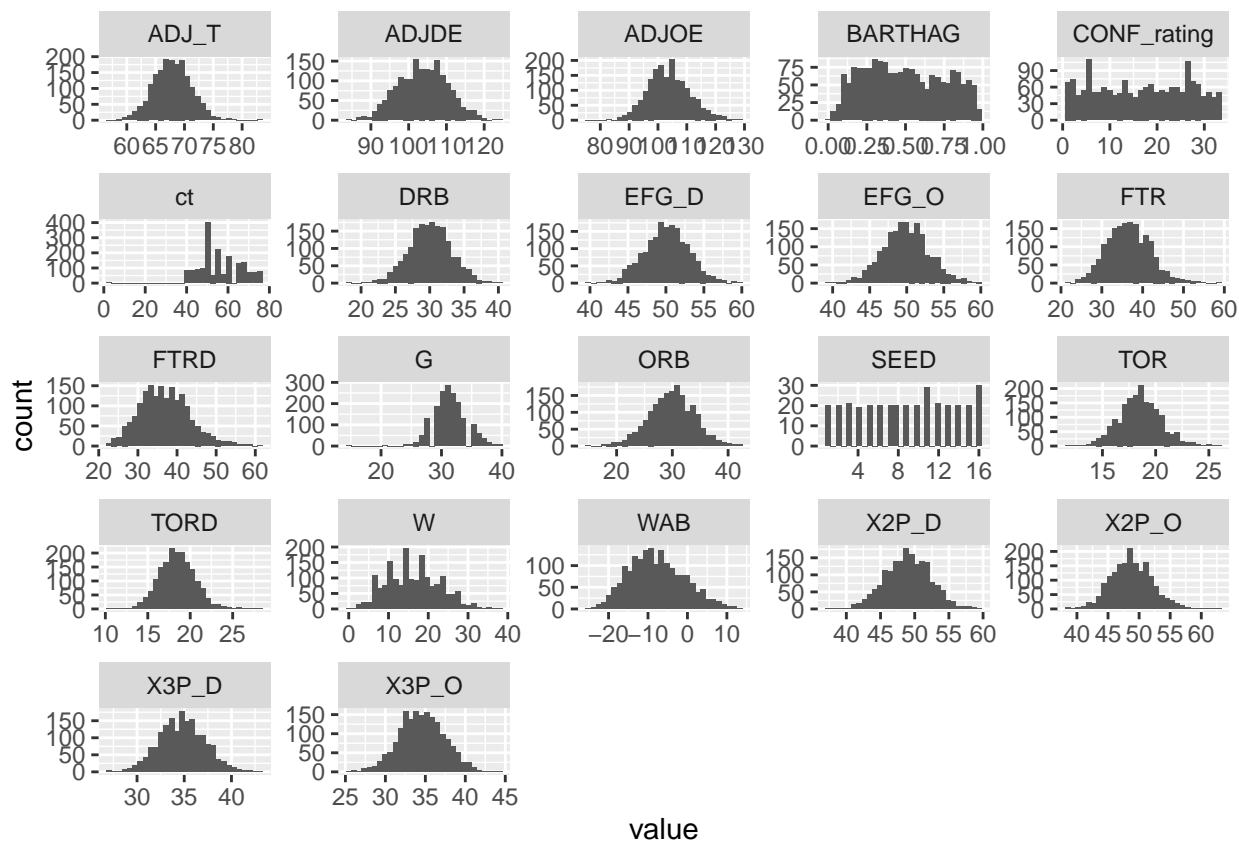
```
## Joining, by = "CONF"
```

```
## Plot a histogram of all the variables to see what the distribution
```

```
data %>% keep(is.numeric) %>% gather() %>% ggplot(aes(value)) + facet_wrap(~ key, scales = "free") + geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 1415 rows containing non-finite values (stat_bin).
```



```
## Convert seed value to binary 1 or 0 response, so that we can have this as our categorical dependent variable
data <- data %>% mutate(SEED = ifelse(data$SEED > 0, 1, 0))

## Same as above, but with the test_data
test_data <- test_data %>% mutate(SEED = ifelse(test_data$SEED > 0, 1, 0))

## Look at the distribution of the y_variable (if a team makes it to the NCAA March Madness Tournament)
data %>% group_by(SEED) %>% summarize(ct = n())
```

```
## # A tibble: 2 x 2
##   SEED    ct
##   <dbl> <int>
## 1     1   340
## 2    NA  1415
```

```
## Convert NA values in SEED column to 0
data[is.na(data)] <- 0
data %>% group_by(SEED) %>% summarize(ct = n())
```

```
## # A tibble: 2 x 2
##   SEED    ct
##   <dbl> <int>
## 1     0  1415
## 2     1   340
```

```
test_data[is.na(test_data)] <- 0
test_data %>% group_by(SEED) %>% summarize(ct = n())
```

```
## # A tibble: 2 x 2
##   SEED   ct
##   <dbl> <int>
## 1     0  285
## 2     1   68
```

```
## Remove TEAM and CONF variable(s) from data set
data <- subset(data, select=-c(Team,CONF))
test_data <- subset(test_data, select=-c(Team,CONF))
```

```
## Look at correlation matrix of variables. Check for multicollinearity between features and little to
cor_matrix <- cor(data)
cor_matrix
```

```
##           G           W      ADJOE      ADJDE      BARTHAG
## G          1.00000000  0.72741530  0.60821747 -0.6083603  0.69111583
## W          0.72741530  1.00000000  0.75481409 -0.7015319  0.82715747
## ADJOE       0.60821747  0.75481409  1.00000000 -0.5084621  0.86753300
## ADJDE      -0.60836034 -0.70153195 -0.50846210  1.00000000 -0.84572219
## BARTHAG     0.69111583  0.82715747  0.86753300 -0.8457222  1.00000000
## EFG_O       0.34162805  0.61287960  0.73018854 -0.2133903  0.54296282
## EFG_D      -0.48457007 -0.60496068 -0.31786887  0.7895983 -0.62287269
## TOR        -0.33984937 -0.46324822 -0.61325708  0.2141322 -0.47963312
## TORD        0.04872899  0.13059830 -0.13307952 -0.2272126  0.03767243
## ORB         0.26286748  0.29884952  0.26458412 -0.2900252  0.31211379
## DRB        -0.19828609 -0.38393679 -0.27348907  0.3762657 -0.35856088
## FTR         0.08875274  0.14307770  0.10068912 -0.1118690  0.12975523
## FTRD       -0.27299135 -0.33633504 -0.35256428  0.2043428 -0.32364957
## X2P_O       0.33278525  0.58075380  0.64739786 -0.2389901  0.51038591
## X2P_D      -0.44042776 -0.53817756 -0.31741363  0.7377599 -0.59242809
## X3P_O       0.23376291  0.44041517  0.58230979 -0.1100500  0.39830254
## X3P_D      -0.36049984 -0.47116430 -0.18708797  0.5542103 -0.41980766
## ADJ_T      -0.04730856 -0.01447517  0.05865379  0.2157609 -0.08131731
## WAB         0.66042635  0.91047203  0.84743224 -0.7992955  0.94064751
## SEED        0.51273772  0.61427628  0.54927685 -0.4947986  0.57569763
## ct          0.25249579  0.14859110  0.21777907 -0.2223562  0.25938951
## CONF_rating -0.22646604 -0.14860795 -0.24703048  0.2257733 -0.26860431
##           EFG_O      EFG_D      TOR      TORD      ORB
## G          0.34162805 -0.48457007 -0.33984937  0.04872899  0.262867476
## W          0.61287960 -0.60496068 -0.46324822  0.13059830  0.298849522
## ADJOE       0.73018854 -0.31786887 -0.61325708 -0.13307952  0.264584124
## ADJDE      -0.21339030  0.78959830  0.21413216 -0.22721259 -0.290025215
## BARTHAG     0.54296282 -0.62287269 -0.47963312  0.03767243  0.312113792
## EFG_O       1.00000000 -0.10441756 -0.36119178 -0.13481262 -0.150245393
## EFG_D      -0.10441756  1.00000000  0.09312867 -0.00253898 -0.351246798
## TOR        -0.36119178  0.09312867  1.00000000  0.10275940  0.105882548
## TORD       -0.13481262 -0.00253898  0.10275940  1.00000000  0.090427807
## ORB        -0.15024539 -0.35124680  0.10588255  0.09042781  1.000000000
## DRB        -0.31717413  0.18169747  0.17302351  0.25245828  0.006822504
```


## FTR	-0.06522234	-0.21155100	0.12741818	0.07369014	0.305153216
## FTRD	-0.37582375	0.11041193	0.28269571	0.35094602	0.129249506
## X2P_0	0.89583705	-0.13607217	-0.29297426	-0.07386592	-0.089119420
## X2P_D	-0.09653918	0.91423521	0.08652803	0.04364090	-0.343974802
## X3P_0	0.77075159	-0.03473390	-0.31614358	-0.16778084	-0.144829053
## X3P_D	-0.07959945	0.72124288	0.05951326	-0.09735751	-0.216605655
## ADJ_T	0.11925478	0.28252025	-0.09095158	-0.03991964	-0.104923300
## WAB	0.56499047	-0.61885255	-0.47827622	0.07869446	0.330585400
## SEED	0.36668569	-0.38009418	-0.30757065	0.07886072	0.223222217
## ct	0.01828070	-0.13930976	-0.08559411	0.03046872	0.167313583
## CONF_rating	-0.11829201	0.10919006	0.16726977	0.07176024	-0.032608267
##	DRB	FTR	FTRD	X2P_0	X2P_D
## G	-0.198286088	0.08875274	-0.27299135	0.33278525	-0.44042776
## W	-0.383936786	0.14307770	-0.33633504	0.58075380	-0.53817756
## ADJOE	-0.273489067	0.10068912	-0.35256428	0.64739786	-0.31741363
## ADJDE	0.376265691	-0.11186896	0.20434279	-0.23899007	0.73775990
## BARTHAG	-0.358560878	0.12975523	-0.32364957	0.51038591	-0.59242809
## EFG_0	-0.317174126	-0.06522234	-0.37582375	0.89583705	-0.09653918
## EFG_D	0.181697467	-0.21155100	0.11041193	-0.13607217	0.91423521
## TOR	0.173023509	0.12741818	0.28269571	-0.29297426	0.08652803
## TORD	0.252458284	0.07369014	0.35094602	-0.07386592	0.04364090
## ORB	0.006822504	0.30515322	0.12924951	-0.08911942	-0.34397480
## DRB	1.000000000	0.09358617	0.25978076	-0.28458627	0.21142781
## FTR	0.093586172	1.000000000	0.24562548	-0.01277812	-0.18990468
## FTRD	0.259780765	0.24562548	1.000000000	-0.35378281	0.10541539
## X2P_0	-0.284586271	-0.01277812	-0.35378281	1.000000000	-0.12008941
## X2P_D	0.211427806	-0.18990468	0.10541539	-0.12008941	1.000000000
## X3P_0	-0.249603179	-0.09109361	-0.26343959	0.41895111	-0.04056623
## X3P_D	0.065672214	-0.14087387	0.08795645	-0.11217027	0.38619936
## ADJ_T	0.005635813	-0.03110244	-0.02897766	0.15341478	0.27760537
## WAB	-0.323179142	0.17169680	-0.32691902	0.53093879	-0.57335587
## SEED	-0.163452376	0.10553267	-0.22081746	0.34025810	-0.33914280
## ct	0.063139199	0.08030849	-0.07650024	0.05873517	-0.13014637
## CONF_rating	0.106462800	0.07739226	0.21496195	-0.11470394	0.11325549
##	X3P_0	X3P_D	ADJ_T	WAB	SEED
## G	0.23376291	-0.36049984	-0.047308561	0.66042635	0.51273772
## W	0.44041517	-0.47116430	-0.014475165	0.91047203	0.61427628
## ADJOE	0.58230979	-0.18708797	0.058653792	0.84743224	0.54927685
## ADJDE	-0.11005002	0.55421026	0.215760876	-0.79929553	-0.49479862
## BARTHAG	0.39830254	-0.41980766	-0.081317307	0.94064751	0.57569763
## EFG_0	0.77075159	-0.07959945	0.119254784	0.56499047	0.36668569
## EFG_D	-0.03473390	0.72124288	0.282520253	-0.61885255	-0.38009418
## TOR	-0.31614358	0.05951326	-0.090951575	-0.47827622	-0.30757065
## TORD	-0.16778084	-0.09735751	-0.039919639	0.07869446	0.07886072
## ORB	-0.14482905	-0.21660566	-0.104923300	0.33058540	0.22322222
## DRB	-0.24960318	0.06567221	0.005635813	-0.32317914	-0.16345238
## FTR	-0.09109361	-0.14087387	-0.031102436	0.17169680	0.10553267
## FTRD	-0.26343959	0.08795645	-0.028977657	-0.32691902	-0.22081746
## X2P_0	0.41895111	-0.11217027	0.153414782	0.53093879	0.34025810
## X2P_D	-0.04056623	0.38619936	0.277605369	-0.57335587	-0.33914280
## X3P_0	1.000000000	-0.01307475	0.029215806	0.41611471	0.27614625
## X3P_D	-0.01307475	1.000000000	0.169875574	-0.44071107	-0.28936929
## ADJ_T	0.02921581	0.16987557	1.000000000	-0.06114658	-0.02533463
## WAB	0.41611471	-0.44071107	-0.061146583	1.000000000	0.64495850

```
## SEED      0.27614625 -0.28936929 -0.025334629  0.64495850  1.00000000
## ct        -0.04334642 -0.10433557 -0.003012734  0.23199150  0.11404568
## CONF_rating -0.07647019  0.06361994  0.047068267 -0.23008867 -0.16459425
##          ct CONF_rating
## G          0.252495789 -0.22646604
## W          0.148591097 -0.14860795
## ADJOE       0.217779071 -0.24703048
## ADJDE      -0.222356217  0.22577328
## BARTHAG     0.259389512 -0.26860431
## EFG_O       0.018280701 -0.11829201
## EFG_D      -0.139309764  0.10919006
## TOR        -0.085594108  0.16726977
## TORD        0.030468715  0.07176024
## ORB         0.167313583 -0.03260827
## DRB         0.063139199  0.10646280
## FTR         0.080308489  0.07739226
## FTRD       -0.076500245  0.21496195
## X2P_O       0.058735172 -0.11470394
## X2P_D      -0.130146367  0.11325549
## X3P_O      -0.043346420 -0.07647019
## X3P_D      -0.104335568  0.06361994
## ADJ_T      -0.003012734  0.04706827
## WAB         0.231991502 -0.23008867
## SEED        0.114045679 -0.16459425
## ct          1.000000000 -0.19638948
## CONF_rating -0.196389482  1.00000000
```

Analysis/Interpretation

```
## First I decided to look at all of the variables to see how many NA values were present.
## After further exploration, I was able to tell the the SEED and POSTSEASON columns were the only vari
## This is because if a team does not make the March Madness tournament, they are not given a postseason
## We view the class type of the data and convert the conference variable to a categorical variable. We
## After plotting a histogram of all of the variables, I can conclude that all variables (except SEED,
## This could cause a potential issue later on because if we do not have a balanced data set in terms of
## I noticed that we had 1,132 rows of NA values in the SEED column. In order to make this a binary pre
## Finally, we remove any unnecessary columns and run a correlation matrix to ensure all variables have
```

```
## Split the data into 10% test set and 90% train set
```

```
### Seeing that the data set is not too large (as far as big data goes), we will have the largest possible
set.seed(42)
```

```
test_index <- createDataPartition(y = data$SEED, times=1, p=0.1, list=FALSE)
```

```
test_set <- data[test_index,]
```

```
train_set <- data[-test_index,]
```

K-fold Cross validation

```
cv_param <- trainControl(method="cv", number = 11)
```

Model Building

Logistic Regression

```
log_reg <- train(SEED~G+ADJOE+ADJDE+ct, data = train_set, method = 'glm')
```

```
## Warning in train.default(x, y, weights = w, ...): You are trying to do
## regression and your outcome only has two possible values Are you trying to do
## classification? If so, use a 2 level factor as your outcome column.
```

```
summary(log_reg)
```


```
##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.72564  -0.21978  -0.05676   0.15542   1.05386
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.0604518  0.2743861  -3.865 0.000116 ***
## G           0.0289815  0.0041222   7.031 3.06e-12 ***
## ADJOE       0.0180153  0.0013780  13.074 < 2e-16 ***
## ADJDE      -0.0133845  0.0015459  -8.658 < 2e-16 ***
## ct          -0.0023993  0.0008356  -2.871 0.004141 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.09697998)
##
##      Null deviance: 249.75  on 1578  degrees of freedom
## Residual deviance: 152.65  on 1574  degrees of freedom
## AIC: 803.8
##
## Number of Fisher Scoring iterations: 2
```

```
log_predictions <- predict(log_reg,test_set)
confusionMatrix(round(log_predictions,digits=0), test_set$SEED)
```

```
##      0  1
## 0 144  3
## 1  14 15
```

```
plotROC(test_set$SEED,log_predictions)
```

ROC Curve



Our logistic regression shows an Area under the curve score of .9253.

Decision Tree

```
tree_ml <- train(SEED~., data = train_set, method = 'rpart', trControl = cv_param)
```

```
## Warning in train.default(x, y, weights = w, ...): You are trying to do
## regression and your outcome only has two possible values Are you trying to do
## classification? If so, use a 2 level factor as your outcome column.
```


```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.
```

```
tree_predictions <- predict(tree_ml, test_set)
confusionMatrix(round(tree_predictions,digits=0), test_set$SEED)
```

```
##      0  1
## 0 146  1
## 1   10 19
```

```
plotROC(test_set$SEED,tree_predictions)
```

ROC Curve



```
### Our decision tree shows an Area under the curve score of .7583.
```

```
## Random Forest
```

```
set.seed(42)
```

```
rf_model <- randomForest(SEED~., data = train_set, boosting=TRUE, trControl = cv_param)
```

```
## Warning in randomForest.default(m, y, ...): The response has five or fewer
```

```
## unique values. Are you sure you want to do regression?
```

```
rf_predictions <- predict(rf_model, test_set)
```


```
confusionMatrix(round(rf_predictions,digits=0),test_set$SEED)
```

```
##      0  1
```

```
## 0 146  1
```

```
## 1  10 19
```

```
plotROC(test_set$SEED,rf_predictions)
```



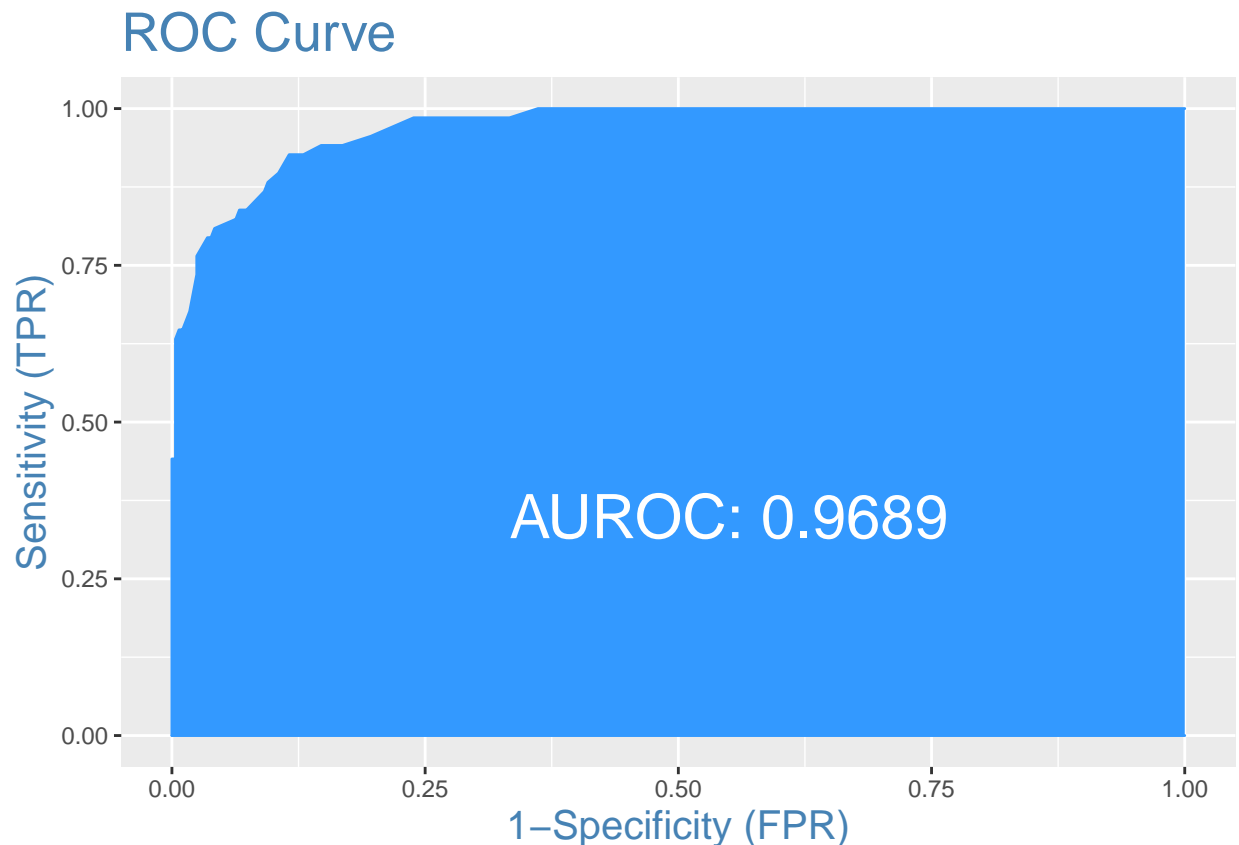
Our random forest shows an Area under the curve score of .935.

Results Using our Best Model

```
final_predictions <- predict(rf_model,test_data)
confusionMatrix(round(final_predictions,digits=0),test_data$SEED)
```

```
##      0  1
## 0 278  7
## 1  18 50
```

```
plotROC(test_data$SEED,final_predictions)
```



```
## We plot an ROC curve to see how much better our model is to someone randomly guessing.
## An ROC Curve plots the True Positive rate (predicted 1 that were true 1's) over our False Positive Rate
## Using our Random Forest model, we are able to predict which teams will make the NCAA March Madness Tournament
## It appears that the most significant variables in predicting whether a team will make the NCAA March Madness Tournament
## Other variables that showed statistical significance include BARTHAG and WAB, but had to be removed due to multicollinearity
## We are also able to see that we had a 89.29% True Positive Rate and a 93.94% True Negative Rate.
## Given the disproportionate number of teams that did not make the tournament, it makes sense that we achieved a high True Negative Rate

# Conclusion
## Given that we had a training dataset of 1,750 teams over 5 years, we could have a slightly overfit model
## Even though this dataset accounts for roughly 55,000+ NCAA basketball games, more data would definitely help
## In conclusion, it appears that we did a great job cleaning up the dataset to predict whether or not a team will make the tournament
## It is currently limited because there are many other factors that could come into play as to why a team will make the tournament
## Some of those factors include: winning conference championships, location of conference championship game, etc.
## These limitations will mostly impact the lower seeds in the bracket as they do not have a guaranteed bid
## In the future, I could have looked into adding more data and incorporating additional factors that would help predict whether a team will make the tournament
```